

A Hierarchical Neural Autoencoder for Paragraphs and Documents

Akilesh B, CS13B1042

Indian Institute of Technology, Hyderabad

November 26, 2016

Introduction

Experiments and
Evaluations

Results

Our
Implementation

- ▶ In this work, we explore an important step toward this generation task: training an LSTM (Long Short Term Memory) auto-encoder to preserve and reconstruct multi-sentence paragraphs.
- ▶ We focus on the component task of training a paragraph (document)-to-paragraph (document) autoencoder to reconstruct the input text sequence from a compressed vector representation from a deep learning model
- ▶ We develop hierarchical LSTM models that arranges tokens, sentences and paragraphs in a hierarchical structure, with different levels of LSTMs capturing compositionality at the token-to-token and sentence-to-sentence levels.

- ▶ Recent LSTM models (Hochreiter and Schmidhuber, 1997) have shown powerful results on generating meaningful and grammatical sentences in sequence generation tasks like machine translation (Sutskever et al., 2014; Bahdanau et al., 2014; Luong et al., 2015) or parsing (Vinyals et al., 2014).
- ▶ This performance is at least partially attributable to the ability of these systems to capture local compositionally: the way neighboring words are combined semantically and syntactically to form meanings that they wish to express.

- ▶ Could these models be extended to deal with generation of larger structures like paragraphs or even entire documents?
- ▶ In standard sequence to sequence generation tasks, an input sequence is mapped to a vector embedding that represents the sequence, and then to an output string of word. Multi-text generation tasks like summarization works in a similar way.
- ▶ Just as the local semantics of words can be captured by LSTM models, can the semantics of higher-level text units (e.g., clauses, sentences, paragraphs, and documents) be captured in a similar way.

Introduction

Experiments and
Evaluations

Results

Our
Implementation

LSTM


A Hierarchical
Neural
Autoencoder for
Paragraphs and
Documents

Introduction

Experiments and
Evaluations

Results

Our
Implementation



- ▶ In this paper, the target is to use an encoder-decoder model to regenerate the document.
- ▶ Specifically, a hierarchical LSTM model is used as an encoder, in which the representations of sentences are learned by a LSTM model whose inputs are words.
- ▶ The representation of the document is learned by another LSTM model whose inputs are sentences representations.
- ▶ A normal LSTM is used as a decoder.

Model 1 : Standard LSTM

- ▶ The whole input and output are treated as one sequence of tokens.
- ▶ Following Sutskever et al. (2014) and Bahdanau et al. (2014), we trained an autoencoder that first maps input documents into vector representations from a $LSTM_{encode}$ and then reconstructs inputs by predicting tokens within the document sequentially from a $LSTM_{decode}$.
- ▶ Two separate LSTMs are implemented for encoding and decoding with no sentence structures considered.

Standard Sequence to Sequence Model

A Hierarchical
Neural
Autoencoder for
Paragraphs and
Documents

Introduction

Experiments and
Evaluations

Results

Our
Implementation




Figure 1: Standard Sequence to Sequence Model.

Model 2 : Hierarchical LSTM

Encoder We first obtain representation vectors at the sentence level by putting one layer of LSTM (denoted as $LSTM_{encode}^{word}$) on top of its containing words:

$$h_t^w(enc) = LSTM_{encode}^{word}(e_t^w, h_{t-1}^w(enc)) \quad (1)$$

The vector output at the ending time-step is used to represent the entire sentence as $e_s = h_{end_s}^w$

Hierarchical LSTM contd. ...

To build representation e_D for the current document/paragraph D , another layer of LSTM (denoted as $LSTM_{encode}^{sentence}$) is placed on top of all sentences, computing representations sequentially for each time step:

$$h_t^s(enc) = LSTM_{encode}^{sentence}(e_t^s, h_{t-1}^s(enc)) \quad (2)$$

Representation $e_{end_D}^s$ computed at the final time step is used to represent the entire document: $e_D = h_{end_D}^s$

Thus one LSTM operates at the token level, leading to the acquisition of sentence-level representations that are then used as inputs into the second LSTM that acquires document-level representations, in a hierarchical structure.

Introduction

Experiments and
Evaluations

Results

Our
Implementation

Hierarchical LSTM contd. ...

Decoder As with encoding, the decoding algorithm operates on a hierarchical structure with two layers of LSTMs. LSTM outputs at sentence level for time step t are obtained by:

$$h_t^s(dec) = LSTM_{decode}^{sentence}(e_t^s, h_{t-1}^s(dec)) \quad (3)$$

The initial time step $h_0^s(d) = e_D$, the end-to-end output from the encoding procedure, $h_t^s(d)$ is used as the original input into the $LSTM_{decode}^{word}$ for subsequently predicting tokens within sentence $t + 1$.

$LSTM_{decode}^{word}$ predicts tokens at each position sequentially, the embedding of which is then combined with earlier hidden vectors for the next time step prediction until the end_s token is predicted.

Hierarchical Sequence to Sequence Model

A Hierarchical
Neural
Autoencoder for
Paragraphs and
Documents

Introduction

Experiments and
Evaluations

Results

Our
Implementation




Figure 2: Hierarchical Sequence to Sequence Model.

- ▶ **Hotel Reviews:** The reviews consisting sentences ranging from 50 to 250 words are considered. The vocabulary set consisted of the 25,000 most frequent words. A special UNK token is used to denote all the less frequent tokens. Reviews that consist of more than 2 percent of unknown words are discarded. The training dataset is comprised of 340,000 reviews; the testing set is comprised of 40,000 reviews.
- ▶ **Wikipedia:** Paragraphs are extracted from Wikipedia corpus that meet the aforementioned length requirements. Paragraphs with larger than 4 percent of unknown words are discarded. The training dataset is comprised of roughly 500,000 paragraphs and testing contains roughly 50,000.

Training and Implementation Details

- ▶ LSTM parameters and word embeddings are initialized from a uniform distribution between $[-0.08, 0.08]$.
- ▶ Each LSTM layer consists of 1000 hidden neurons and the dimensionality of word embeddings is set to 1000
- ▶ Decoding algorithm allows generating at most 1.5 times the number of words in inputs.

Evaluation Metric

We use ROUGE-1, ROUGE-2 scores on the dataset.

ROUGE is a recall-oriented measure widely used in the language generation literature. It measures the n-gram recall between the candidate text and the reference (generated) text(s).

$$ROUGE_n = \frac{\sum_{gram_n \in input} count_{match}(gram_n)}{\sum_{gram_n \in input} count(gram_n)} \quad (4)$$

where $count_{match}$ denotes the number of n-grams co-occurring in the input and output.

ROUGE How much the words (and/or n-grams) in the human reference text appeared in the machine generated text. It measures recall.

BLEU Purely measuring recall will inappropriately reward long outputs. BLEU is designed to address such an issue by emphasizing precision. n-gram precision scores for our situation are given by:

$$precision_n = \frac{\sum_{gram_n \in output} count_{match}(gram_n)}{\sum_{gram_n \in output} count(gram_n)} \quad (5)$$

BLEU How much the words (and/or n-grams) in the machine generated text appeared in the human reference text. It measures precision.

Table: Results for the LSTM Models

Model	Dataset	BLEU	ROUGE-1	ROUGE-2
Standard	Hotel Review	0.241	0.571	0.302
Standard	Wikipedia	0.178	0.502	0.228
Hierarchical	Hotel Review	0.267	0.590	0.330
Hierarchical	Wikipedia	0.202	0.529	0.250

Hierarchical LSTM examples

Input-Wiki:

*paris is the capital and most populous city of france .
situated on the seine river , in the north of the country , it is
in the centre of the le-de-france region . the city of paris has
a population of 2273305 inhabitants . this makes it the fifth
largest city in the european union measured by the
population within the city limits .*

Output-Wiki:

*paris is the capital and most populated city in france .
located in the UNK , in the north of the country , it is the
center of UNK . paris , the city has a population of NUM
inhabitants . this makes the eu s population within the city
limits of the fifth largest city in the measurement*

Hierarchical LSTM examples ...

Input-Review:

on every visit to nyc , the hotel beacon is the place we love to stay . so conveniently located to central park , lincoln center and great local restaurants . the rooms are lovely . beds so comfortable , a great little kitchen and new wizz bang coffee maker . the staff are so accommodating and just love walking across the street to the fairway supermarket with every imaginable goodies to eat .

Output-Review

every time in new york , lighthouse hotel is our favorite place to stay . very convenient , central park , lincoln center , and great restaurants . the room is wonderful , very comfortable bed , a kitchenette and a large explosion of coffee maker . the staff is so inclusive , just across the street to walk to the supermarket channel love with all kinds of what to eat .

Introduction

Experiments and
Evaluations

Results

Our
Implementation

The Signal Media One-Million News Articles Dataset is used for training and testing. It contains 1 million articles that are mainly English, but they also include non-English and multi-lingual articles.

Each article in the dataset has the following fields.

- ▶ **id**: a unique identifier for the article.
- ▶ **title**: the title of the article.
- ▶ **content**: the textual content of the article (which occasionally contained HTML and JavaScript content).
- ▶ **source**: the name of the article source (e.g. Reuters).
- ▶ **published**: the publication date of the article.
- ▶ **media-type**: either "News" or "Blog".

Introduction

Experiments and
Evaluations

Results

Our
Implementation

- ▶ Minimal pre-processing such as lower-casing, separating on the basis of punctuation marks are applied.
- ▶ All the punctuations are removed in order to reduce the vocabulary size.
- ▶ We remove the articles with more than 20% label the tokens which occur less than 10 times in our corpus.
- ▶ For every article, cosine similarity of each sentence present in it, is computed with the corresponding title. We retain only the sentences with top 5 cosine similarity scores per article.
- ▶ The average number of words in an article is 122 .

Constructing feature vector and training

- ▶ We initially trained a Word2vec model on our whole corpus, which is a two-layer neural net that generates word embeddings. We found out that the results were not good.
- ▶ We used a one-hot encoding for our feature vector. This means the size of feature vector will be equal to the vocabulary.
- ▶ The standard LSTM encoder decoder model is implemented in Keras / Theano.
- ▶ Adagrad/Adadelata is used as optimizer for training.

Training details

Table: Hyperparameter list

Hyperparameter	Value
Hidden layer size	500/1000
Batch size	128
Learning rate	0.1
Optimizer	Adagrad/Adadelata

- The training was performed on a Tesla K-20 GPU with 6GB RAM.

Standard LSTM examples

Generated Text:

*the new remote control about thicker than ray previous
changed ipad apps of case girl s the good thing*

Original Text:

*the new remote control is thicker than the previous one and
in this case that s a good thing*

Generated Text: (Bad Example)

to we've difference and diversity we've been by

Original Text:

where we've succeeded, its because we've been different.

Hierarchical LSTM examples

Generated Text:

*if we spread malicious gossip we our poisoning words own
peace demoralise sabotaging situations own influence*

Original Text:

*if we spread malicious gossip we are poisoning our own
future and sabotaging our own influence with others*

Generated Text:

*pressing down on the screen starts different actions
depending fathers how much force healing exerted*

Original Text:

*pressing down on the screen starts different actions
depending on how much force is exerted*

Introduction

Experiments and
Evaluations

Results

Our
Implementation

Hierarchical LSTM examples ...

Generated Text:

*the misunderstood tried searching i action movies the
comedy movies the the remote control*

Original Text:

*i also tried searching for action movies and comedy movies
with the remote control s microphone*

Generated Text:

*bowling ball on third 3rd side with ease redzone left texas
went to chris nutall*

Original Text:

*bowling ball on third 3rd and in the redzone texas state went
to chris nutall senior running back*

Table: Results for the LSTM Models

Model	No of articles	BLEU	ROUGE-1	ROUGE-2
Standard	500	0.1267	0.4113	0.2591
Hierarchical	200	0.2372	0.8143	0.6014
Hierarchical	500	0.3513	0.7713	0.5945

Challenges faced

- ▶ Computational constraints - GPU.
- ▶ Huge size of vocabulary due to one-hot encoding.
- ▶ Hyperparameter optimizations which ensure error function converges to an appropriate local minima.

References

A Hierarchical
Neural
Autoencoder for
Paragraphs and
Documents

Introduction

Experiments and
Evaluations

Results

Our
Implementation

 A Hierarchical Neural Autoencoder for Paragraphs and Documents. Jiwei et al.

 Sequence to sequence learning with Neural Networks. Sutskever et al.

 <http://research.signalmedia.co/newsir16/signal-dataset.html>